DR Detection Report

Team: Insulin EYEdentifiers
Abhishikt Mahajan, Karan Naik

27 October 2024

1 Introduction

1.1 Problem Statement

Diabetic retinopathy (DR) is a serious complication of diabetes that can lead to vision loss. This project aims to develop an automated classification system for DR severity using deep learning models, which can assist healthcare professionals in early diagnosis and management.

1.2 Motivation

Early and accurate detection of diabetic retinopathy is essential for preventing vision loss in diabetic patients. The complexity of DR classification stems from the subtle differences between severity levels, requiring precise and sensitive models. A robust automated classification system could drastically improve diagnostic speed and accuracy, particularly in settings with limited access to ophthalmologists. Building such a system required an iterative approach, experimenting with various deep learning architectures, and analyzing the trade-offs in terms of performance, computational efficiency, and complexity. After rigorous testing, EfficientNet and ResNet34 were chosen as the best candidates due to their demonstrated ability to perform well on image classification tasks. This project represents an in-depth exploration of model performance optimization for DR detection, highlighting the process of model selection, fine-tuning, and evaluation.

1.3 Objectives and Work Plan

The primary objectives of this project include:

- **Developing deep learning models (EfficientNet and ResNet34):** Multiple architectures were trained and evaluated to select models that could accurately differentiate DR severity. EfficientNet and ResNet34 were fine-tuned for this purpose.
- **Performance comparison and iterative improvement:** Each model's accuracy, loss, precision, and recall were tracked across training iterations, and hyperparameters were adjusted to maximize performance.
- **Evaluation and testing strategy:** A systematic testing approach was applied, including cross-validation, learning rate scheduling, and regularization, to achieve reliable performance and avoid overfitting.

The work plan involved the following stages:

1. **Data preprocessing and augmentation**: The dataset was cleaned and augmented to increase variability and robustness during training.
2. **Model training**: EfficientNet and ResNet34 models were trained on the processed dataset, with intermediate evaluations to ensure convergence.
3. **Hyperparameter tuning and optimization**: Learning rates, batch sizes, and optimizer parameters were fine-tuned through trials to improve model performance.
4. **Evaluation on validation and test sets**: Both models were evaluated on validation and test sets, comparing performance metrics to ensure generalizability.

5. **Result analysis and final model selection**: The results were analyzed and documented, highlighting EfficientNet as the top-performing model based on validation metrics.

2 Materials and Methods

2.1 System Architecture

The architecture includes preprocessing of retinal images, training EfficientNet and ResNet34 models, and evaluating them on a validation set. Both models were trained using cross-entropy loss and Adam optimizer with learning rate scheduling for optimal convergence.

2.2 Model Architectures

- **EfficientNet**: This model scales depth, width, and resolution, achieving high accuracy with fewer parameters. The final layers are adjusted for multi-class classification.
- **ResNet34**: A residual network that preserves gradient flow through skip connections, minimizing vanishing gradients.

2.3 Flowchart

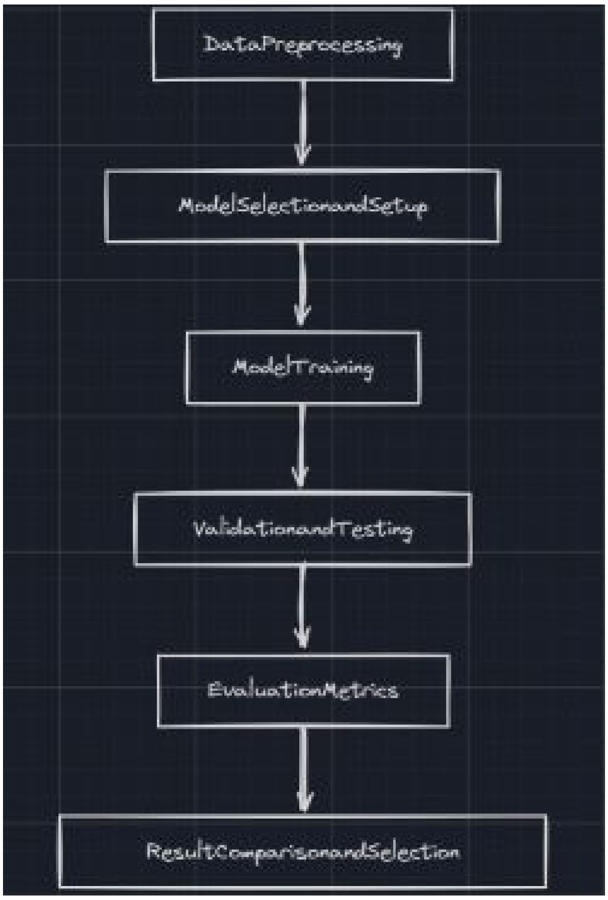

Figure 1: System Flowchart for DR Detection

2.4 Performance Metrics

Models were evaluated based on:

- **Accuracy**: Percentage of correctly classified instances.

- **Loss**: Cross-entropy loss indicating prediction error.
- **Precision, Recall, F1-Score**: Assess balanced performance across DR stages.
- **ROC AUC**: Assesses class discrimination ability, useful in multi-class classification.

3 Results and Discussion

3.1 EfficientNet Results

Figure 2: EfficientNet - Training and Validation Accuracy and Loss vs Epochs

EfficientNet achieved a final validation accuracy of approximately 81%, with loss decreasing over epochs, showing effective convergence.

3.2 ResNet34 Results

Figure 3: ResNet34 - Training and Validation Accuracy and Loss vs Epochs

ResNet34 reached around 80% validation accuracy, showing similar convergence behavior with stable performance.

3.3 Result Comparison

Both models performed well, with EfficientNet slightly outperforming ResNet34. Fine-tuning could further narrow the gap or improve accuracy.

4 Conclusion and Future Work

This project successfully demonstrates the potential of deep learning models, specifically EfficientNet and ResNet34, in automating the classification of diabetic retinopathy (DR) severity. Through rigorous experimentation and hyperparameter tuning, EfficientNet emerged as the top-performing model, achieving a high validation accuracy with efficient computational resource usage. This outcome reinforces the suitability of scalable models like EfficientNet for complex medical image analysis tasks where model performance and resource efficiency are both crucial.

The project has shown that deep learning can substantially aid in the early detection of diabetic retinopathy, potentially enabling faster and more accurate diagnoses. This automated approach is particularly beneficial in resource-limited healthcare environments, where access to trained specialists may be limited.

Future work will focus on enhancing this solution by exploring ensemble methods to leverage strengths across multiple architectures. Additionally, further investigation into data augmentation techniques and model interpretability could help improve robustness and provide clearer insights into model predictions. Integrating these models into real-world healthcare applications, supported by ongoing refinement and testing, has the potential to make a meaningful impact on diabetic retinopathy screening and management.

5 References

- Tan, M., Le, Q. V. (2019). EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. *Proceedings of the International Conference on Machine Learning (ICML).*
- He, K., Zhang, X., Ren, S., Sun, J. (2016). Deep Residual Learning for Image Recognition. *IEEE Conference on Computer Vision and Pattern Recognition (CVPR).*